

Fundamentals: Markov Decision Process

**Kianté Brantley
CS 1840 Fall 2026**

Course Overview

Instructor: Kianté Brantley

TFs: Alex Cai, Philip Nielsen, Egil Rolstad


Course Overview

Instructor: Kianté Brantley

TFs: Alex Cai,

Philip Nielsen,


Egil Rolstad




Course Overview

Instructor: Kianté Brantley


TFs: Alex Cai,



Philip Nielsen,



Egil Rolstad



CAs: Emmanuel Rassou,



Justin Li,



Ishaan Sinha



Announcements

Office Hours

website: <https://harvard-cs-1840.github.io/2026/>

Type	Day	Start	End	Staff	Campus	Location
Office Hours	Monday	3:45 PM	5:30 PM	Kianté	Allston	SEC 5.402
Office Hours	Monday	8:00 PM	9:30 PM	Justin	Cambridge	Mather Pset Nights
Office Hours	Thursday	2:15 PM	4:00 PM	Emmanuel	Allston	SEC 2.122/2.123
Office Hours	Thursday	4:30 PM	6:15 PM	Alex	Cambridge	Cruft Lab 309
Office Hours	Tuesday	3:45 PM	5:30 PM	Egil	Allston	SEC 2.122/2.123
Office Hours	Tuesday	7:00 PM	8:45 PM	Philip	Cambridge	Pierce Hall 301
Office Hours	Wednesday	7:00 PM	8:45 PM	Ishaan	Cambridge	Winthrop CS Nights
Section	Monday	2:15 PM	3:30 PM	Alex	Allston	SEC 1.412
Section	Thursday	12:45 PM	2:00 PM	Emmanuel	Allston	SEC 1.412
Section	Tuesday	11:15 AM	12:30 PM	Egil	Allston	SEC 1.412
Section	Tuesday	4:30 PM	5:45 PM	Philip	Cambridge	Pierce Hall 213 (Brooks Room)
Section	Wednesday	4:30 PM	5:45 PM	Ishaan & Justin	Cambridge	Maxwell Dworkin G135

Announcements

Homework 1 released 9/11 **(this Friday)**

Homework 0 due 9/17 **(next Friday)**

If any part of Homework 0 feels challenging or unfamiliar,
please reach out to the TFs/CAs for help.

Reading Materials

Three Books:

1) Reinforcement Learning: Theory and Algorithms

by Alekh Agarwal, Kianté Brantley, Nan Jiang, Sham Kakade, and Wen Sun.

This is an advanced RL theory book and is mostly out of scope for this class.

Please let us know if you find any typos or mistakes in the book

2) An Introduction to Reinforcement Learning

by Alex Cai

3) Reinforcement Learning: An Introduction

by Richard Sutton and Andrew Barto.

This is the classic text in reinforcement learning.


Questions?

Outline:

1. Definitions of Markov Decision Process
2. Value functions (V and Q functions)
3. Bellman equations


The Mathematical framework:

Infinite horizon Markov Decision Process




The Mathematical framework:

Infinite horizon Markov Decision Process




The Mathematical framework:

Infinite horizon Markov Decision Process



The Mathematical framework:

Infinite horizon Markov Decision Process




$$r(s, a), s' \sim P(\cdot | s, a)$$

$P(\cdot | s, a)$: distribution over the next state


The Mathematical framework:

Infinite horizon Markov Decision Process




The Mathematical framework:


Infinite horizon Markov Decision Process



Example: 2-D simple car navigation




(a)




(b)

Example: 2-D simple car navigation

$$s = [x, y, \theta, v]^T \in \mathbb{R}^4$$



(a)




(b)


Example: 2-D simple car navigation

$$s = [x, y, \theta, v]^T \in \mathbb{R}^4$$

$$a = [\alpha, \phi]^T \in \mathbb{R}^2$$




(a)




(b)

Example: 2-D simple car navigation



(a)




(b)

$$s = [x, y, \theta, v]^T \in \mathbb{R}^4$$


$$a = [\alpha, \phi]^T \in \mathbb{R}^2$$

$$r(s, a) = \begin{cases} 100 & (x, y) \in \mathcal{X}_{\text{goal}} \\ -1 & \text{hit obstacles} \\ 0 & \text{else} \end{cases}$$

Example: 2-D simple car navigation



(a)




(b)

$$s = [x, y, \theta, v]^T \in \mathbb{R}^4$$


$$a = [\alpha, \phi]^T \in \mathbb{R}^2$$

$$r(s, a) = \begin{cases} 100 & (x, y) \in \mathcal{X}_{\text{goal}} \\ -1 & \text{hit obstacles} \\ 0 & \text{else} \end{cases}$$
$$s' = f(s, a) + \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, I)$$

Example: 2-D simple car navigation



(a)



(b)

$$s = [x, y, \theta, v]^{\top} \in \mathbb{R}^4$$

$$a = [\alpha, \phi]^{\top} \in \mathbb{R}^2$$


$$r(s, a) = \begin{cases} 100 & (x, y) \in \mathcal{X}_{\text{goal}} \\ -1 & \text{hit obstacles} \\ 0 & \text{else} \end{cases}$$

$$s' = f(s, a) + \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, I)$$

$$f(s, a) = \begin{bmatrix} x + \tau v \cos \theta \\ y + \tau v \sin \theta \\ \theta + \tau v \tan(\phi)/d \\ v + \tau \alpha \end{bmatrix}$$


Example: Robotic Pick-and-Place

robot hand needs to pick the ball and hold it in a goal (x,y,z) position



Example: Robotic Pick-and-Place


robot hand needs to pick the ball and hold it in a goal (x,y,z) position



State s : robot configuration (e.g., joint angles)
and the ball's position

Example: Robotic Pick-and-Place

robot hand needs to pick the ball and hold it in a goal (x,y,z) position




State s : robot configuration (e.g., joint angles) and the ball's position

Action a : Torque on joints in arm & fingers

Example: Robotic Pick-and-Place

robot hand needs to pick the ball and hold it in a goal (x,y,z) position




State s : robot configuration (e.g., joint angles) and the ball's position

Action a : Torque on joints in arm & fingers

Transition $s' \sim P(\cdot | s, a)$: physics + some noise

Example: Robotic Pick-and-Place

robot hand needs to pick the ball and hold it in a goal (x,y,z) position




State s : robot configuration (e.g., joint angles) and the ball's position

Action a : Torque on joints in arm & fingers


Transition $s' \sim P(\cdot | s, a)$: physics + some noise

Cost $c(s, a)$: torque magnitude + dist to goal

Example: OpenAI Gym demonstrations




Example: OpenAI Gym demonstrations




Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]




Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]



Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]



$$r(s, a) = \begin{cases} 1 & \text{pole angle} \in [-12^\circ, 12^\circ], \\ 0 & \text{else} \end{cases}$$

Policy

$$\mathcal{S} \rightarrow \Delta(\mathcal{A})$$

A mapping from state to action (what action should I take if I'm in this state...)

$$a \sim \pi(\cdot | s)$$

Policy

$$\mathcal{S} \rightarrow \Delta(\mathcal{A})$$

A mapping from state to action (what action should I take if I'm in this state...)

$$a \sim \pi(\cdot | s)$$

conditional
distribution over all
actions

Policy

$$\mathcal{S} \rightarrow \Delta(\mathcal{A})$$

A mapping from state to action (what action should I take if I'm in this state...)

$$a \sim \pi(\cdot | s)$$

conditional
distribution over all
actions

Deterministic vs stochastic?

Policy

$$\mathcal{S} \rightarrow \Delta(\mathcal{A})$$

A mapping from state to action (what action should I take if I'm in this state...)

$$a \sim \pi(\cdot | s)$$


conditional
distribution over all
actions

Deterministic vs stochastic?

Q: Assume S state and A actions, how many different deterministic policies we can have?

Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]




$$r(s, a) = \begin{cases} 1 & \text{pole angle} \in [-12^\circ, 12^\circ], \\ 0 & \text{else} \end{cases}$$

Policy 1: uniform random
 $\pi(0 | s) = \pi(1 | s) = 0.5, \forall s$

Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]



$$r(s, a) = \begin{cases} 1 & \text{pole angle} \in [-12^\circ, 12^\circ], \\ 0 & \text{else} \end{cases}$$


Policy 1: uniform random
 $\pi(0 | s) = \pi(1 | s) = 0.5, \forall s$

Policy 2: adaptive

$$\pi(s) = \begin{cases} 0(\text{left}) & \text{if pole angle} < 0 \\ 1(\text{right}) & \text{else} \end{cases}$$

Example: OpenAI Gym demonstrations

State = [cart pos, cart velocity,
pole angle, pole angular velocity]



$$r(s, a) = \begin{cases} 1 & \text{pole angle} \in [-12^\circ, 12^\circ], \\ 0 & \text{else} \end{cases}$$

Policy 1: uniform random
 $\pi(0 | s) = \pi(1 | s) = 0.5, \forall s$

Policy 2: adaptive

$$\pi(s) = \begin{cases} 0(\text{left}) & \text{if pole angle} < 0 \\ 1(\text{right}) & \text{else} \end{cases}$$

Outlines:

- ✓ 1. Definitions of Markov Decision Process
- 2. Value functions (V and Q functions)
- 3. Bellman equations

Performance of a policy

Expected total reward of a policy π :

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, \pi \right]$$

$\gamma \in [0,1)$: discount factor (value future reward less and less)

Performance of a policy

Expected total reward of a policy π :

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, \pi \right]$$

$\gamma \in [0,1)$: discount factor (value future reward less and less)

Q: think about the CartPole example, is there a way we can estimate $V^\pi(s)$ at a given s ?

Optimal policy

$\pi^\star$: the policy that maximizes expected future reward at all states

$$V^\star(s) \geq V^\pi(s), \forall s, \forall \pi$$

Optimal policy

$\pi^\star$: the policy that maximizes expected future reward at all states

$$V^\star(s) \geq V^\pi(s), \forall s, \forall \pi$$

Fact: $\pi^\star$ is deterministic, only depends on state (i.e., Markovian);
 $\pi^\star$ always exists for any infinite horizon discounted MDP

Optimal policy

$\pi^\star$: the policy that maximizes expected future reward at all states

$$V^\star(s) \geq V^\pi(s), \forall s, \forall \pi$$

Fact: $\pi^\star$ is deterministic, only depends on state (i.e., Markovian);
 $\pi^\star$ always exists for any infinite horizon discounted MDP

Q: what is the optimal policy when $\gamma = 0$?

State-action Q function

$$Q^{\pi}(s, a) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, a_0 = a, \pi \right]$$

Outlines:

✓ 1. Definitions of Markov Decision Process

✓ 2. Value functions (V and Q functions)


3. Bellman equations

Can we express the value function using only one step?

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, a \sim \pi \right]$$


Can we express the value function using only one step?

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, a \sim \pi \right]$$




Can we express the value function using only one step?

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, a \sim \pi \right]$$



Can we express the value function using only one step?

$$V^\pi(s) = \mathbb{E} \left[r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^h r_h + \dots \mid s_0 = s, a \sim \pi \right]$$



$$V^\pi(s) = \mathbb{E}_{a \sim \pi(\cdot|s)} \left[r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} [V^\pi(s')] \right]$$

Bellman equation for value function

Can we express the value function using only one step?

Your homework: understand the one-step relationship between V and Q

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot | s, a)} [V^\pi(s')]$$

$$V^\pi(s) = \mathbb{E}_{a \sim \pi(\cdot | s)} Q^\pi(s, a)$$

Summary:

Discounted infinite horizon MDP:

State, action, policy, transition, reward (or cost), discount factor

V function and Q function

Key concept: **Bellman equation**

Next Lecture

- Exact Policy Evaluation
- Fixpoint solution
- Approximate Policy Evaluation via an Iterative Algorithm

Summary

Feedback

<https://bit.ly/4gYn2nC>

